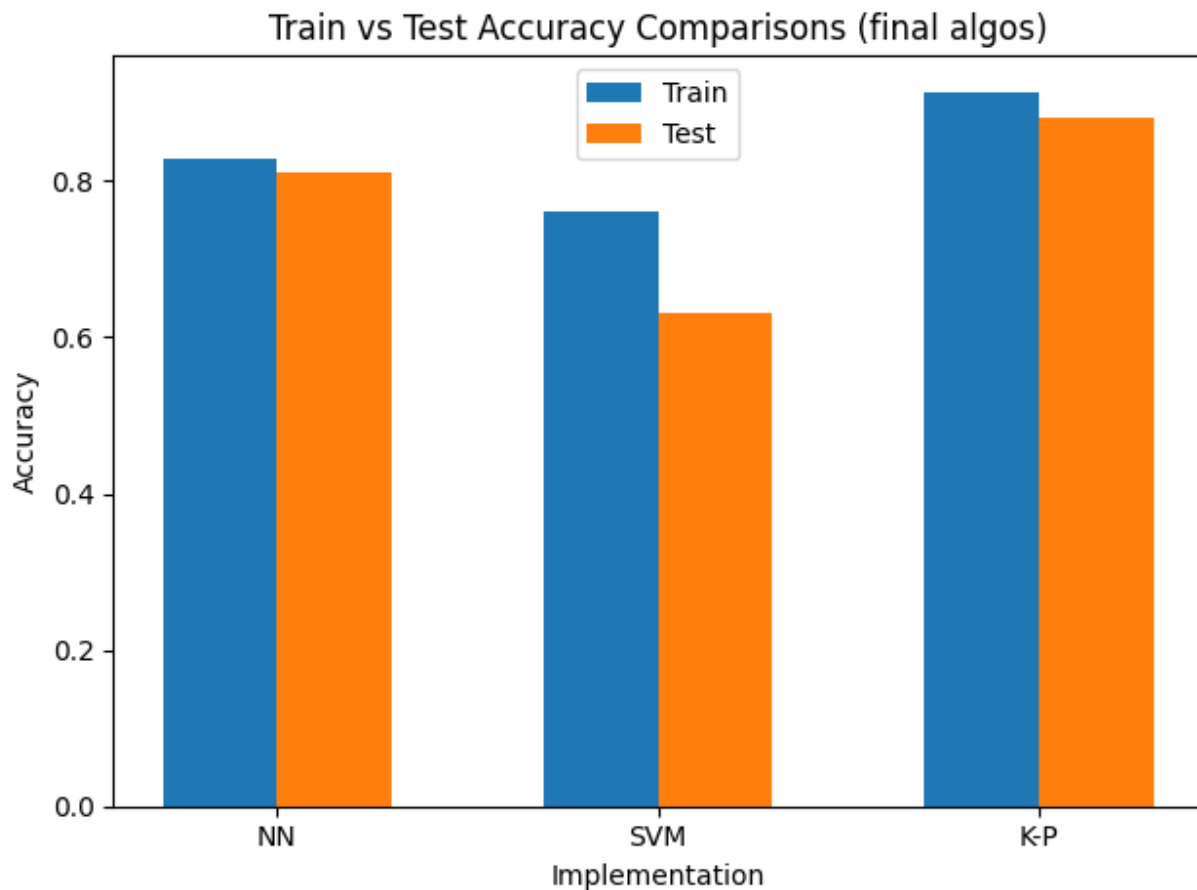
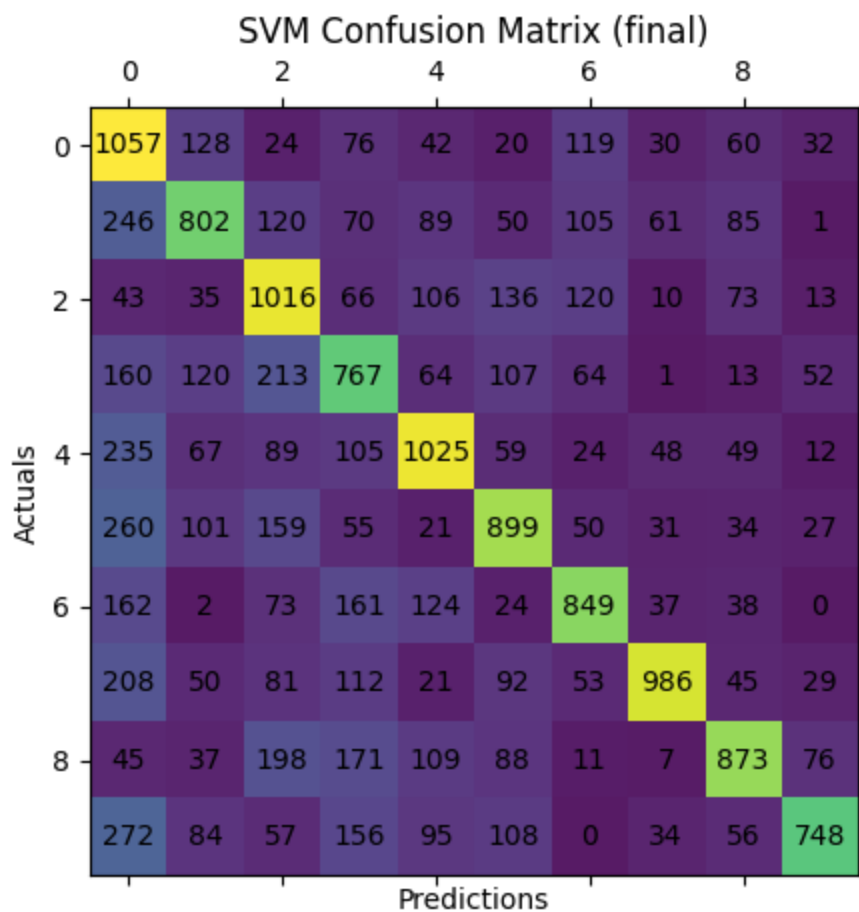




Mini-Project 2

Included are train vs test accuracies for each method, along with a confusion matrix against the test data. Below that is code from the relevant files for the dataset handler, final neural network, and SVM/kernel perceptron with ECOC.

- Perceptron performed the best at 0.914/0.881 train/test accuracy after a lot of tweaking. It's biggest opportunity from the confusion matrix seems to be classifying 4's, which it confuses with 9's relatively frequently. This is understandable due to the similarities between the digits.
- Neural network was the runner up at 0.829/0.812 train/test accuracy. A large portion of the inaccuracies come from misclassifying 9's. The network chooses a 4, or to a lesser extent a 0, more often than it correctly classifies 9's. It is curious that we do not see similar misclassifications of the 4's, but perhaps the model ran away with leaning towards outputting a 4 when in doubt, as digits 5 through 8 also see a relatively high level of misclassification as a 4.
- Finally, SVM struggled the most at 0.761/0.631 train/test accuracies. I poked and prodded, but could not achieve a better performance as some absurd fit times were slowing me down. The confusion matrix doesn't give much insight, besides that the method seems to fall back on outputting 0 or 2 more than anything else when in doubt.
- Had difficulty finding a happy medium of iterations throughout. Especially in the neural network, it was a fine line between not enough iterations to learn meaningful information and so many that the model predicted the training data at 100% but failed to generalize well. I'm thinking this may be a large opportunity for the SVM method, as it was difficult to run trials with the fit times as they were.







```
import numpy as np
from sklearn.datasets import load_svmlight_file
import pickle

class DataSet(object):
    def __init__(self):

        self.data, self.labels = self.readPickles()
        self.testData, self.testLabels = [], []
        self.trainData, self.trainLabels = [], []
        self.inputSize = len(self.data[0])
        self.hiddenSize = 100
        self.outputSize = 10

        self.hiddenValues = []
        self.outputValues = []
        self.inputWeights = []
        self.outputWeights = []
        self.biases = np.random.rand(2,1)

    # must have initialized dataset with openData()
    def readPickles(self):
        with open('data.scale', 'rb') as f:
            data = pickle.load(f)
            f.close()
```

```

        with open('label.scale', 'rb') as f:
            label = pickle.load(f)
        f.close()
        return data, label

    def openData(self):
        x, label = load_svmlight_file('MnistORIG.scale')
        x = x.toarray()
        label = label.astype(int)
        return x, label

    def setLabels(self):
        oldLabels = self.labels
        self.labels = np.zeros((len(oldLabels), 10))
        for i in range(len(oldLabels)):
            index = oldLabels[i]
            self.labels[i][index] = 1

    def splitData(self, splitAt):
        self.testData, self.testLabels = self.data[splitAt:],
        self.labels[splitAt:]
        self.trainData, self.trainLabels = self.data[:splitAt],
        self.labels[:splitAt]

import numpy as np
import operator
import time
from manageDataset import DataSet
from sklearn.metrics import confusion_matrix
import pickle

def sig(x, deriv=False):
    if (deriv==True):
        return sig(x) * (1 - sig(x))
    return 1/(1+np.exp(-x))

def train():
    startTime = time.time()
    X = log.trainData
    y = log.trainLabels

    np.random.seed(1)

    web1 = 2 * np.random.random((780,100)) - 1
    web2 = 2 * np.random.random((100,10)) - 1
    bias = 2 * np.random.random((2,1)) - 1
    np.seterr(all='ignore')
    for i in range(1000):
        inputChain = X
        hiddenChain = sig(np.dot(inputChain, web1) + bias[0][0])
        outputChain = sig(np.dot(hiddenChain, web2) + bias[1][0])

        outputError = y - outputChain
        if (i%100) == 0:
            print(f"Iteration {i} Error: {float(np.mean(np.abs(outputError)))}")

```

```

        outputChange = outputError * sig(outputChain, deriv=True)
        hiddenError = outputChange.dot(web2.T)
        hiddenChange = hiddenError * sig(hiddenChain, deriv=True)
        biasChange1 = np.mean(hiddenError) * sig(bias[0][0], deriv=True)
        biasChange2 = np.mean(outputError) * sig(bias[1][0], deriv=True)

        web2 += hiddenChain.T.dot(outputChange)
        web1 += inputChain.T.dot(hiddenChange)
        bias[0][0] += biasChange1
        bias[1][0] += biasChange2
    log.inputWeights = web1
    log.outputWeights = web2
    print(f"Finished Training, took {time.time() - startTime} seconds")

def predict(predData, trueLabels):
    i = 0
    correct = 0
    realLabels = []
    predLabels = []
    for item in predData:
        log.hiddenValues = sig(np.dot(item, log.inputWeights))
        log.outputValues = sig(np.dot(log.hiddenValues, log.outputWeights))
        index, value = max(enumerate(log.outputValues),
key=operator.itemgetter(1))
        indexTrue, valueTrue = max(enumerate(trueLabels[i]),
key=operator.itemgetter(1))
        predLabels.append(index)
        realLabels.append(indexTrue)
        #print(f"Prediction: {index} vs Actual: {indexTrue}")
        if index == indexTrue:
            correct += 1
        i += 1
    print(f"Accuracy: {correct/len(predData)}")
    confMatrix = confusion_matrix(realLabels, predLabels)
    print(confMatrix)
    with open('conf.txt', 'wb') as f:
        pickle.dump(confMatrix, f)
    f.close()

if __name__ == "__main__":
    log = DataSet()
    log.setLabels()
    log.splitData(42000)
    print("Training...")
    train()
    print("Accuracy against Training Data")
    predict(log.trainData, log.trainLabels)
    print("Accuracy against Test Data")
    predict(log.testData, log.testLabels)

import numpy as np
from manageDataset import DataSet

def binaryTestTrim(log):
    xVals = []

```

```

yVals = []
for i in range(len(log.labels)):
    if log.labels[i] == 0:
        yVals.append(-1)
        xVals.append(log.data[i])
    elif log.labels[i] == 1:
        yVals.append(1)
        xVals.append(log.data[i])
    else:
        pass
log.data = xVals
log.labels = yVals

def linearSVM(log, iters):
    weights = pegSVM(log.trainData, log.trainLabels, iters)
    errors = 0
    for i in range(len(log.testLabels)):
        decision = weights @ log.testData[i].T
        if decision < 0:
            prediction = -1
        else:
            prediction = 1
        if prediction != log.testLabels[i]: errors += 1
    return 1 - errors/len(log.testLabels)

def kernFunc(x, y):
    mean = np.linalg.norm(x - y)**2
    variance = 1
    return np.exp(-mean/(2*variance))

def multiSVM(log, iters):
    weights = pegPercKernel(log.trainData, log.trainLabels, kernFunc, iters)
    errors = 0
    for i in range(len(log.testLabels)):
        decision = 0
        for j in range(len(log.testLabels)):
            decision += weights[j] * log.trainLabels[j] *
kernFunc(log.trainData[j], log.testData[i])
        if decision < 0:
            prediction = -1
        else:
            prediction = 1
        if prediction != log.testLabels[i]: errors += 1
    return 1 - errors/len(log.testLabels)

def pegSVM(x, y, iterations, lam=0.1):
    weights = np.zeros(x[0].shape)
    for i in range(iterations):
        iterCount = randint(0, len(y)-1)
        step = 1/(lam*(i+1))
        decision = y[iterCount] * weights @ x[iterCount].T
        if decision < 1:
            weights = (1 - step*lam) * weights + step*y[iterCount]*x[iterCount]
        else:
            weights = (1 - step*lam) * weights
    return weights

```

```

def pegPercKernel(x, y, kernel, iterations, lam=0.1):
    weights = np.zeros(len(y))
    for _ in range(iterations):
        print(f"Peg Iteration: {_}")
        it = randint(0, len(y)-1)
        decision = 0
        for j in range(len(y)):
            decision += weights[j] * y[it] * kernel(x[it], x[j])
        decision *= y[it]/lam
        if decision < 1:
            weights[it] += 1
    return weights

def testBase():
    log = DataSet()
    binaryTestTrim(log)
    log.splitData(42000)
    iters = 1000
    kernel = False

    accuracy = multiSVM(log, iters)
    print('Accuracy Against Test:', accuracy)

def sepData(log, id):
    _log = {'data': [], 'labels': []}

    for i in range(len(log.trainLabels)):
        _log['data'].append(log.trainData[i])
        if log.trainLabels[i] == id:
            _log['labels'].append(1)
        else:
            _log['labels'].append(-1)
    _log['data'] = np.array(_log['data'])
    _log['labels'] = np.array(_log['labels'])
    return _log

def testECOC():
    log = DataSet()
    log.splitData(7000)
    iters = 1000
    weightClasses = []
    for i in range(log.outputSize):
        print("using RBf")
        _log = sepData(log, i)
        weightClasses.append(pegPercKernel(_log['data'], _log['labels'],
kernFunc, iters))
    errors = 0
    for i in range(len(log.testLabels)):
        predictions = []
        for k in range(10):
            weights = weightClasses[k]
            decision = 0
            for j in range(len(log.trainLabels)):
                decision += weights[j] * log.trainLabels[j] *
kernFunc(log.trainData[j], log.testData[i])
            predictions.append(decision)

```

```

        predictions = np.array(predictions)
        classLabels = predictions.argmax()
        if classLabels != log.testLabels[i]:
            errors += 1
            print(f"Error: {classLabels} found, {log.testLabels[i]} expected,
index {i}")
        accuracy = 1 - errors / len(log.testLabels)
        print(f"Error: {errors / len(log.testLabels)}")
        print(f"Accuracy: {accuracy}")

```

```

if __name__ == "__main__":
    testBase()
    testECOC()

```

```

import numpy as np
from manageDataset import DataSet

```

```

class PolynomialPerceptron():
    def __init__(self, inputs, targets, n, p, max_iter):
        self.inputs = inputs
        self.targets = targets
        self.p = p
        self.n = n
        self.max_iter = max_iter
        self.alpha = np.zeros(n)

    def polyFunc(self, X, Y, p):
        Y = Y.T
        res = np.dot(X, Y)
        res = np.add(res, Y)
        res = np.exp(res, p)
        return res

    def train(self):
        K = self.polyFunc(self.inputs, self.inputs, self.p)
        for i in range(self.max_iter):
            print(f"Iteration {i}")
            for j in range(self.n):
                alph = self.alpha * self.targets
                k = alph * K[j]
                if self.targets[j] * k <= 0:
                    self.alpha[j] += 1

    def predict(self, inputData):
        tmp = (self.alpha * self.targets) * inputData
        testPred = 1 if tmp > 0 else 0
        return testPred

    def acc(self, inputs, targets):
        correct = 0
        kVal = self.polyFunc(inputs, self.inputs, self.p)
        for idx, each in enumerate(inputs):
            correct += self.predict(kVal[idx]) == targets[idx]
        return correct / len(inputs)

```

```

log = DataSet()
log.splitData(42000)

```



```

iters = 1000
n = 42000
c = 0.01

model = PolynomialPerceptron(log.trainData, log.trainLabels,n,c,iters)
model.train()
acc = model.acc(log.testData, log.testLabels)
print(f"accuracy: {acc}")
confMat = confusion_matrix(log.realLabels, log.testLabels)
with open('percConf.txt', 'wb') as f:
    pickle.dump(confMat, f)
f.close()

```